

# MACHINE LEARNING CLASSIFICATION ASSIGNMENT REPORT & CODE - SOMESH GAUR

```
In [158]: ### THIS CODE IS REQUIRED AS I AM WORKING ON IBM WATSON PLATFORM ###

# @hidden_cell
# The project token is an authorization token that is used to access project resources like data sources, connections, and used by platform APIs.
from project_lib import Project
project = Project(project_id='00018ec1-6194-427c-8ce8-67d29662e2a8', project_access_token='p-7e008f5b3a3bbf2681434775bc94dd48755a4fec')
pc = project.project_context
# Write "Wine_Quality_Data.csv" to working directory
open('/home/wsuser/work/Wine_Quality_Data.csv', 'wb').write(project.get_file('Wine_Quality_Data.csv').read())
```

Out[158]: 464197

```
In [160]: import pandas as pd
import numpy as np
#import os -> remove if not required
import seaborn as sns
import matplotlib.pyplot as plt
#import plotly.express as px -> remove if not required
wdata=pd.read_csv('/home/wsuser/work/Wine_Quality_Data.csv', sep=',')
```

## A) Main objective of the analysis and the benefits that analysis provides to the business or stakeholders of this data.

**MAIN OBJECTIVE:** The main objective of this exercise is to build and evaluate various Classification models to predict whether a particular wine is "good quality" or not based on their features. While it may be interesting to understand the interpretability of the model, but from practical and usefulness stand-point, I would like to focus on predictability of model in this exercise.

**BENEFITS:** While many people like wines but very few people can really differentiate between quality of wines. It may be of interest and importance to individual, restaurants and wineries to know what makes a good wine. But such prediction by a human might vary from person to person and require highly skilled tasters. It will be very useful if we can develop machine learning models to predict a good quality wine based on its physical/chemical composition.

## **B) Brief description of the data set, a summary of its attributes, and an outline of objective of analysis.**

BRIEF DESCRIPTION OF DATASET: I am using Wine\_Quality\_Data.csv which a popular dataset used for ML models. This data set contains various chemical properties of wine, such as acidity, sugar, pH, and alcohol. It also contains a quality metric (3-9, with highest being better) and a color (red or white).

This dataset is originally based on two datasets that are related to red and white variants of the Portuguese "Vinho Verde" wine. Due to privacy and logistic issues, only physicochemical (inputs) and sensory (the output) variables are available (e.g. there is no data about grape types, wine brand, wine selling price, etc.).

### **SUMMARY OF ATTRIBUTES :**

There are 13 features and 6497 rows in the dataset. Input variables (based on physicochemical tests):

1 - fixed acidity , 2 - volatile acidity , 3 - citric acid , 4 - residual sugar, 5 - chlorides

6 - free sulfur dioxide , 7 - total sulfur dioxide, 8 - density , 9 - pH , 10 - sulphates ,

11 - alcohol , 12- Color (Physical quality) , 13 - Quality (score between 0 and 10) - Output variable (based on sensory data)

BRIEF DESCRIPTION of 12 variables and 1 output variable (quality).

Fixed Acidity: are non-volatile acids that do not evaporate readily

Volatile Acidity: are high acetic acid in wine which leads to an unpleasant vinegar taste

Citric Acid: acts as a preservative to increase acidity. When in small quantities, adds freshness and flavor to wines

Residual Sugar: is the amount of sugar remaining after fermentation stops. The key is to have a perfect balance between sweetness and sourness. It is important to note that wines > 45g/ltrs are sweet

Chlorides: the amount of salt in the wine

Free Sulfur Dioxide: it prevents microbial growth and the oxidation of wine

Total Sulfur Dioxide: is the amount of free + bound forms of SO<sub>2</sub>

Density: sweeter wines have a higher density

pH: describes the level of acidity on a scale of 0–14. Most wines are always between 3–4 on the pH scale

Alcohol: available in small quantities in wines makes the drinkers sociable

Sulphates: a wine additive that contributes to SO<sub>2</sub> levels and acts as an antimicrobial and antioxidant

Quality: which is the output variable/predictor

Color : Color of wine (Red/White)

Datatypes: There are float64 :11 , int64 :1 and object: 1

OBJECTIVE OF ANALYSIS : The main aim of this exercise is to predict quality of wine based on physiochemical features. I am going to develop various classification models and evaluate which of them is best in predicting the quality of wine based on the above 12 features.

```
In [4]: print(wdata.dtypes)
```

```
fixed_acidity      float64
volatile_acidity   float64
citric_acid        float64
residual_sugar     float64
chlorides          float64
free_sulfur_dioxide float64
total_sulfur_dioxide float64
density            float64
pH                 float64
sulphates          float64
alcohol            float64
quality            int64
color              object
dtype: object
```

```
In [4]: wdata.dtypes.value_counts()
```

```
Out[4]: float64    11
        int64      1
        object     1
        dtype: int64
```

```
In [3]: wdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed_acidity          6497 non-null   float64
1   volatile_acidity       6497 non-null   float64
2   citric_acid            6497 non-null   float64
3   residual_sugar         6497 non-null   float64
4   chlorides              6497 non-null   float64
5   free_sulfur_dioxide    6497 non-null   float64
6   total_sulfur_dioxide   6497 non-null   float64
7   density               6497 non-null   float64
8   pH                    6497 non-null   float64
9   sulphates             6497 non-null   float64
10  alcohol               6497 non-null   float64
11  quality               6497 non-null   int64
12  color                 6497 non-null   object
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB
```

```
In [5]: wdata.shape
```

```
Out[5]: (6497, 13)
```

## C) Brief summary of data exploration and actions taken for data cleaning and feature engineering.

```
In [9]: ### WRTIE SUMMARY ###

#DATA EXPLORATION :

#DATA CLEANING :

#FEATURE ENGINEERING :
```

### C-1 : DATA EXPLORATION

In [4]: `wdata.head()`

Out[4]:

	<b>fixed_acidity</b>	<b>volatile_acidity</b>	<b>citric_acid</b>	<b>residual_sugar</b>	<b>chlorides</b>	<b>free_sulfur_dioxide</b>	<b>total_s</b>
<b>0</b>	7.4	0.70	0.00	1.9	0.076	11.0	
<b>1</b>	7.8	0.88	0.00	2.6	0.098	25.0	
<b>2</b>	7.8	0.76	0.04	2.3	0.092	15.0	
<b>3</b>	11.2	0.28	0.56	1.9	0.075	17.0	
<b>4</b>	7.4	0.70	0.00	1.9	0.076	11.0	


In [7]: `wdata.describe().T`

Out[7]:

	<b>count</b>	<b>mean</b>	<b>std</b>	<b>min</b>	<b>25%</b>	<b>50%</b>	<b>75%</b>	
<b>fixed_acidity</b>	6497.0	7.215307	1.296434	3.80000	6.40000	7.00000	7.70000	15.
<b>volatile_acidity</b>	6497.0	0.339666	0.164636	0.08000	0.23000	0.29000	0.40000	1.
<b>citric_acid</b>	6497.0	0.318633	0.145318	0.00000	0.25000	0.31000	0.39000	1.
<b>residual_sugar</b>	6497.0	5.443235	4.757804	0.60000	1.80000	3.00000	8.10000	65.
<b>chlorides</b>	6497.0	0.056034	0.035034	0.00900	0.03800	0.04700	0.06500	0.
<b>free_sulfur_dioxide</b>	6497.0	30.525319	17.749400	1.00000	17.00000	29.00000	41.00000	289.
<b>total_sulfur_dioxide</b>	6497.0	115.744574	56.521855	6.00000	77.00000	118.00000	156.00000	440.
<b>density</b>	6497.0	0.994697	0.002999	0.98711	0.99234	0.99489	0.99699	1.
<b>pH</b>	6497.0	3.218501	0.160787	2.72000	3.11000	3.21000	3.32000	4.
<b>sulphates</b>	6497.0	0.531268	0.148806	0.22000	0.43000	0.51000	0.60000	2.
<b>alcohol</b>	6497.0	10.491801	1.192712	8.00000	9.50000	10.30000	11.30000	14.
<b>quality</b>	6497.0	5.818378	0.873255	3.00000	5.00000	6.00000	6.00000	9.


```
In [8]: # Data distribution of wine color/type.  
sns.countplot(x='color',data=wdata)  
wdata['color'].value_counts()  
### Clearly, we have data imbalance here with much more data with respect to r  
ed wine ###
```

```
Out[8]: white    4898  
red      1599  
Name: color, dtype: int64
```




```
In [9]: # Let's check data distribution of wine quality, which is our target variable
sns.countplot(x='quality',data=wdata)
wdata.quality.value_counts()
### Clearly, we have data imbalance here with much more data with respect to Q
uality=5,6,7 ###
```

```
Out[9]: 6    2836
        5    2138
        7    1079
        4     216
        8     193
        3      30
        9       5
Name: quality, dtype: int64
```



```
In [15]: plt.figure(figsize=(6,4))
sns.countplot(x='color', hue='quality',data=wdata)
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x7fad6021a610>
```

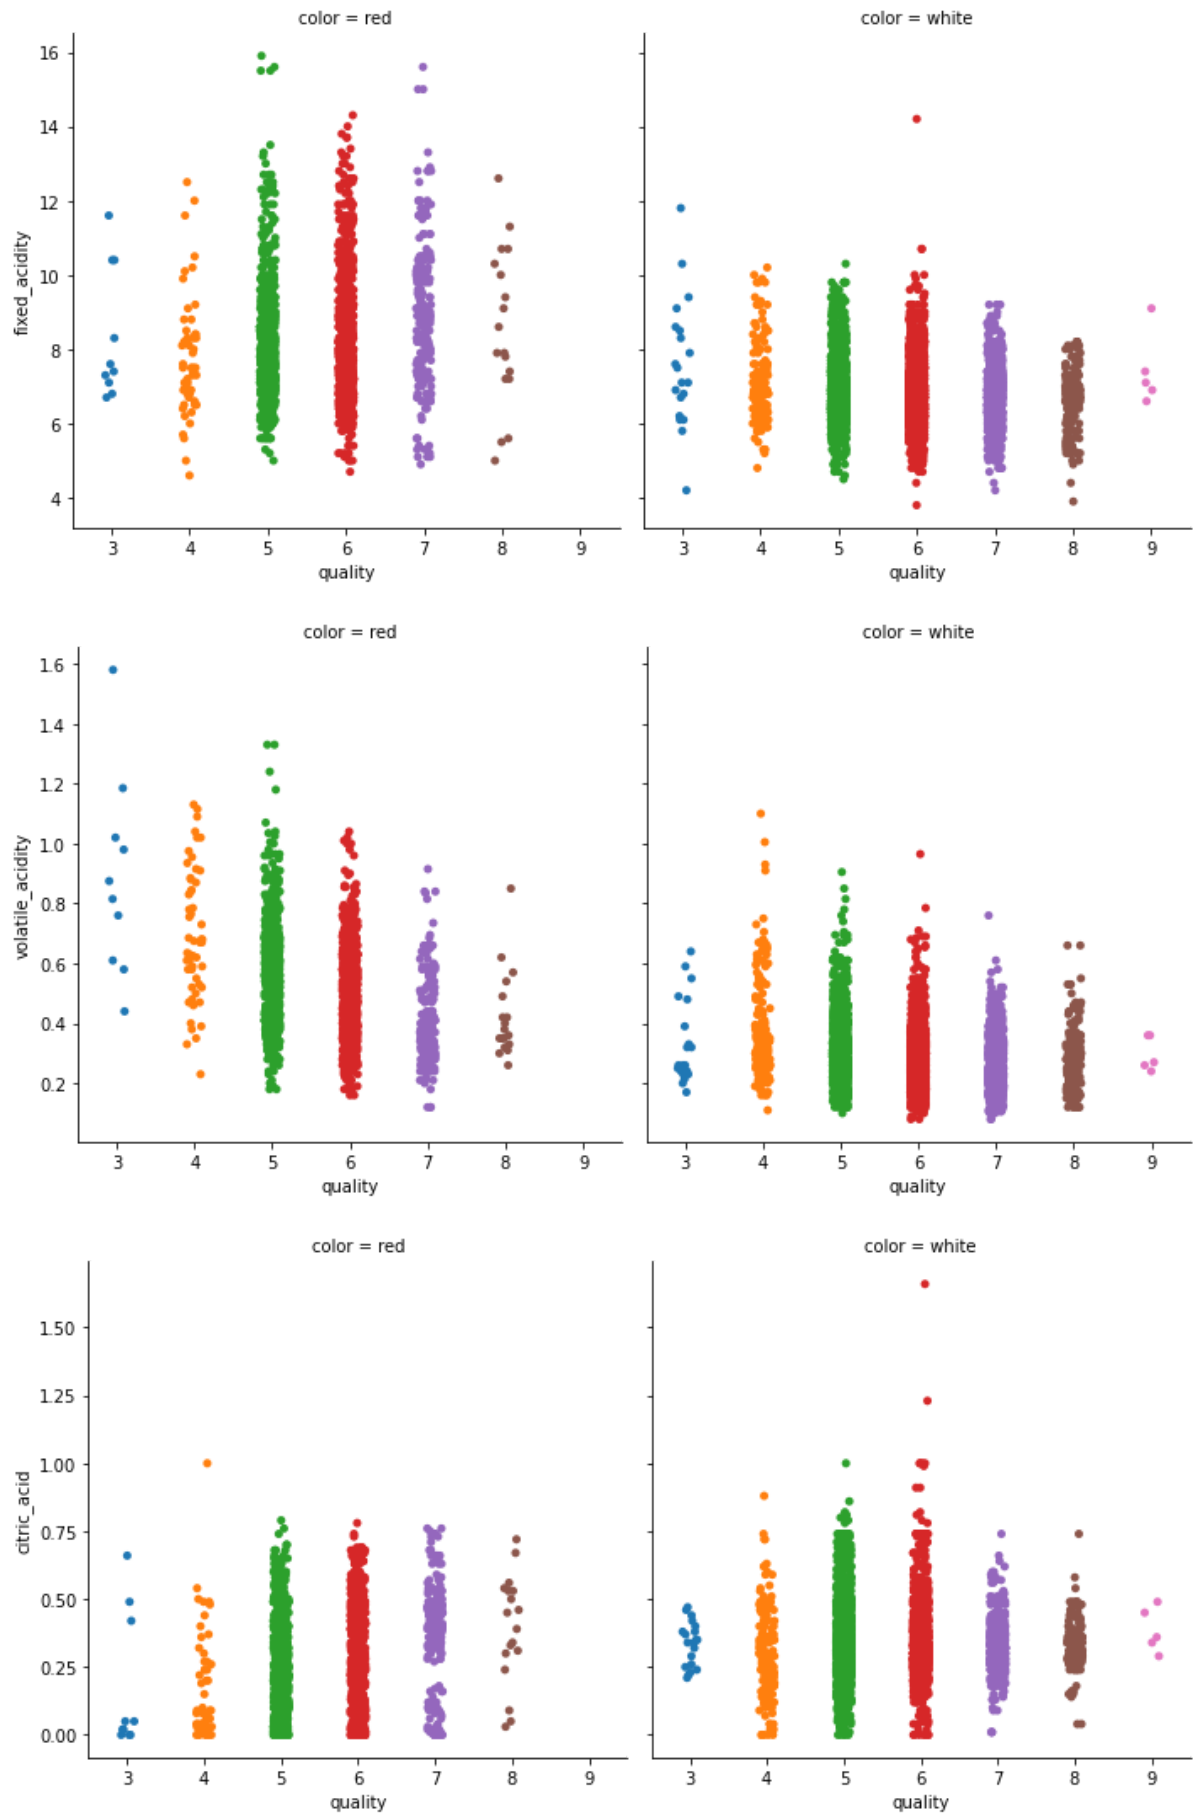


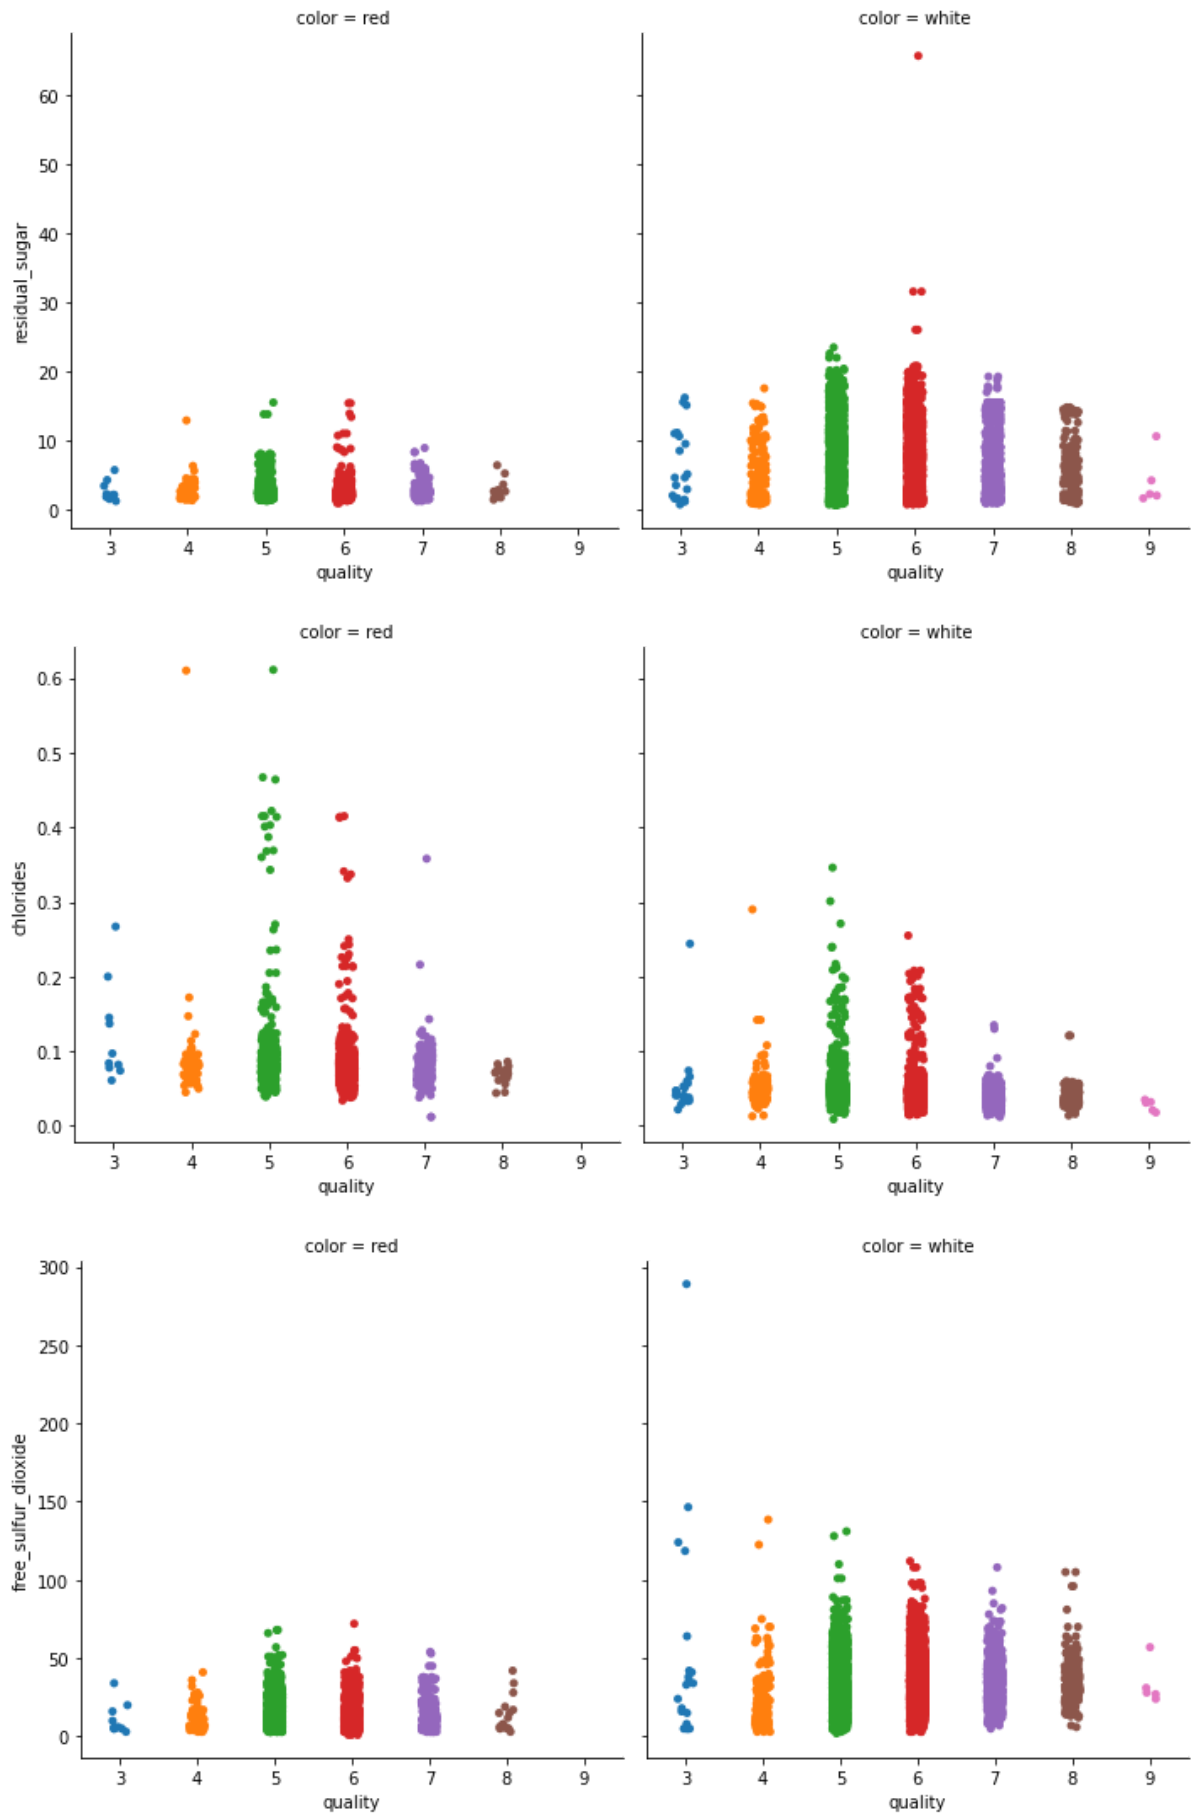
```
In [54]: # numerical columns
numeric_data = [feature for feature in wdata.columns if wdata[feature].dtypes
               != 'O' and feature not in 'quality']

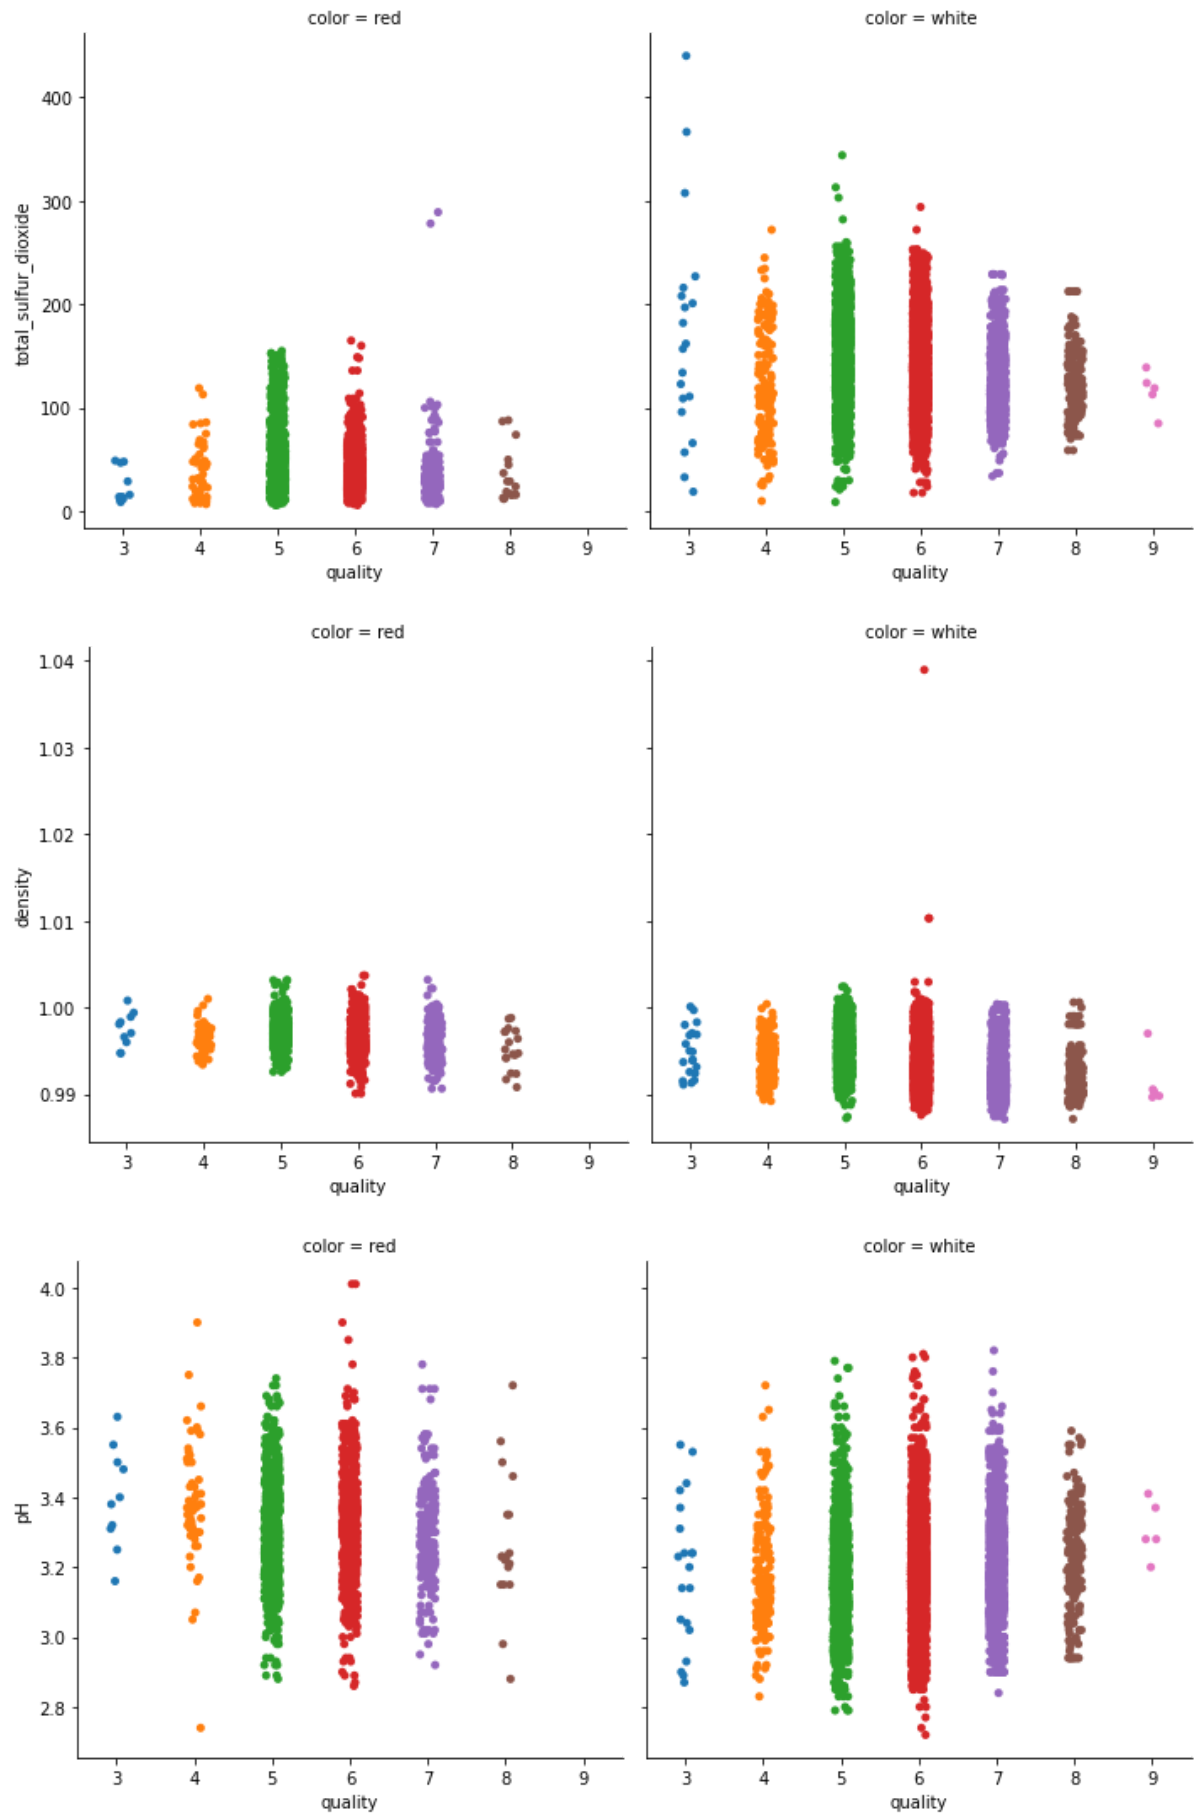
# Let's see the distribution of data with Quality ( on red and white wine )
# Also, we can develop some sense of 'Suspected outliers'

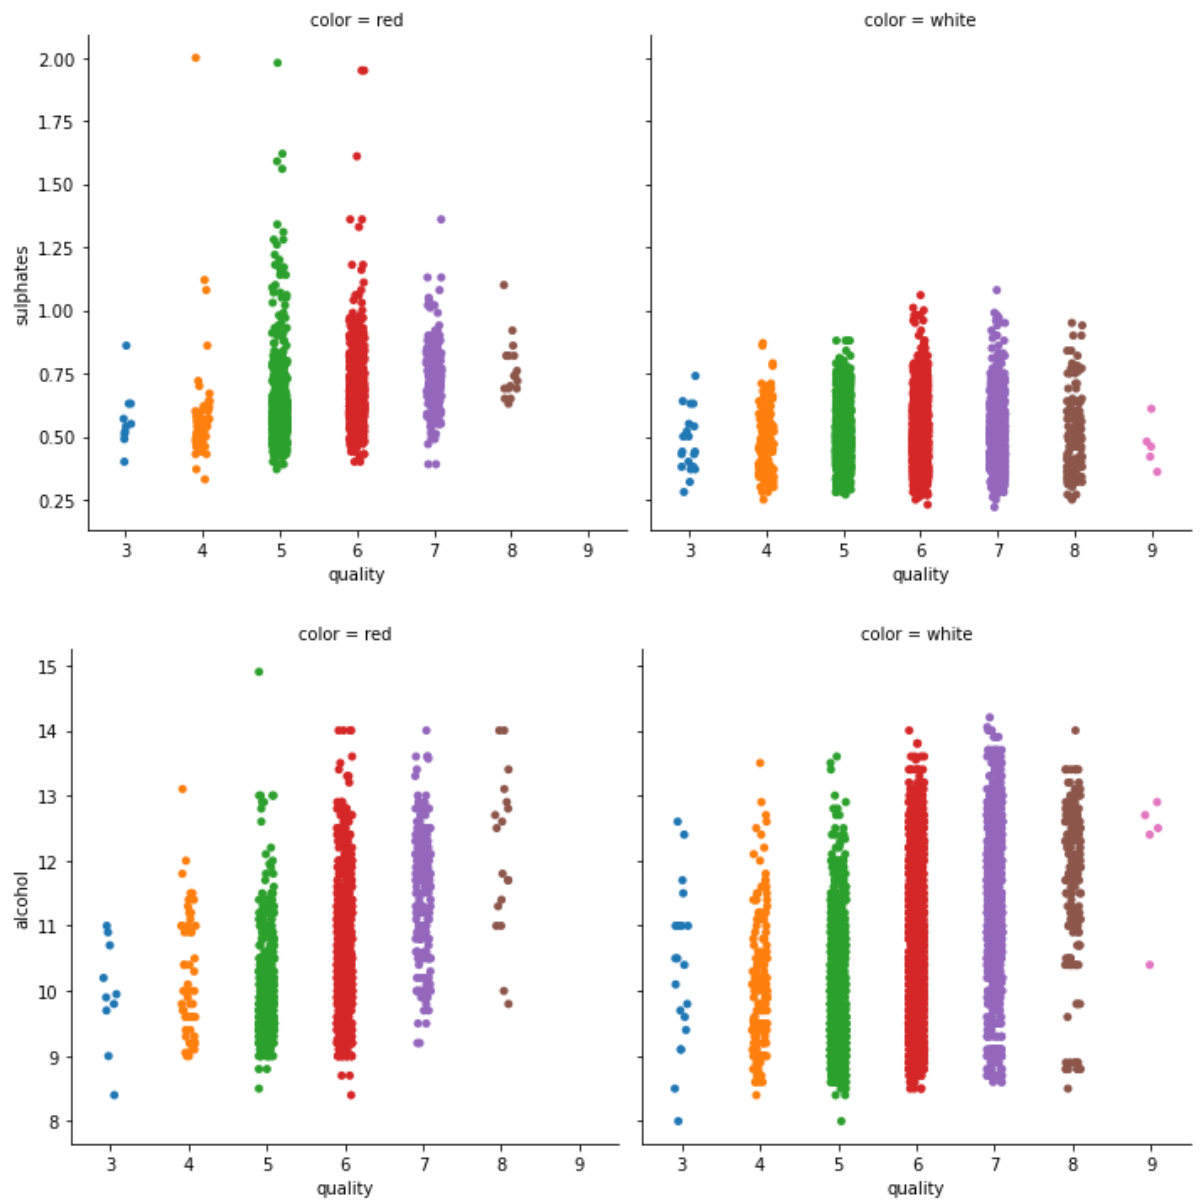
for feature in numeric_data:
    sns.catplot(x='quality',y=feature,col='color',data=wdata)
    plt.show()
```




















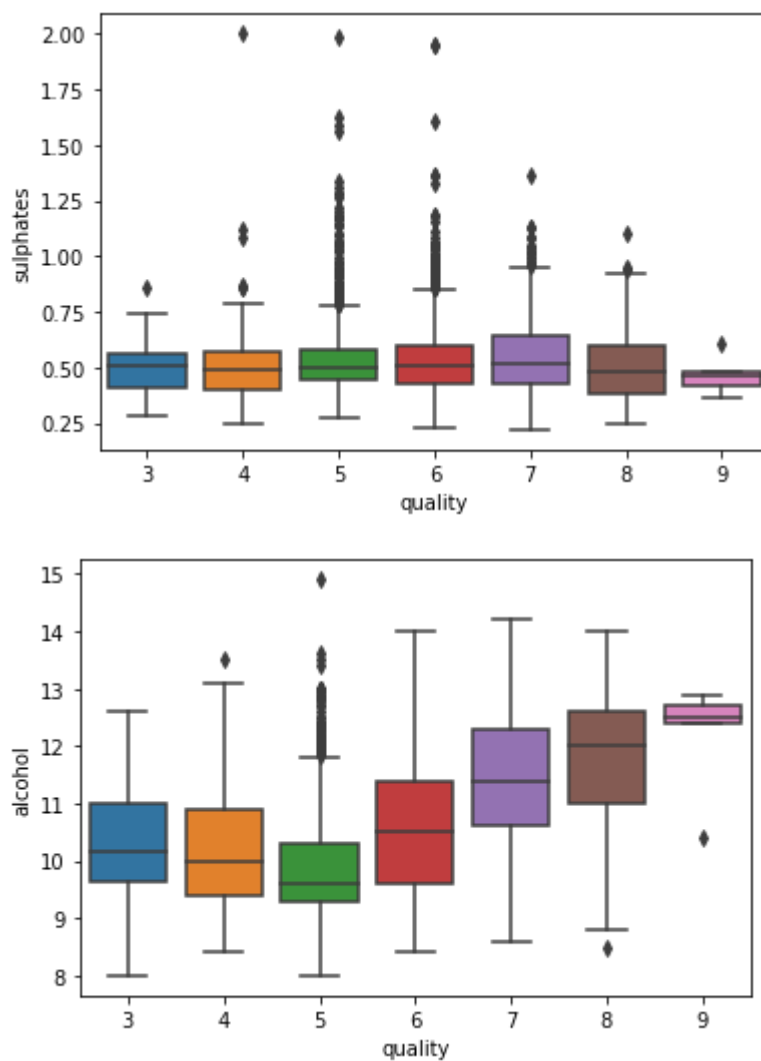
```
In [55]: for feature in numeric_data:
          sns.boxplot(x='quality',y=feature,data=wdata)
          plt.show()
```











```
In [32]: sns.set_context('talk')  
#sns.set_palette(palette)  
sns.set_style('white')
```

```
In [34]: sns.pairplot(wdata, hue='color')
```

/opt/conda/envs/Python-3.7-OpenCE/lib/python3.7/site-packages/seaborn/distributions.py:369: UserWarning: Default bandwidth for data is 0; skipping density estimation.

```
warnings.warn(msg, UserWarning)
```

/opt/conda/envs/Python-3.7-OpenCE/lib/python3.7/site-packages/seaborn/distributions.py:369: UserWarning: Default bandwidth for data is 0; skipping density estimation.

```
warnings.warn(msg, UserWarning)
```


```
Out[34]: <seaborn.axisgrid.PairGrid at 0x7f0abd314390>
```



```
In [11]: # Correlation Matrix
# Will check the correlations between the variables to get a much better understanding of the relationships between variables
##### there are some variables that are strongly correlated to quality.
### These variables are likely to be most important features in our machine Learning model


corr = wdata.corr()
plt.subplots(figsize=(15,10))
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns, annot=True, cmap=sns.diverging_palette(220, 20, as_cmap=True))
```

Out[11]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7fad61e697d0>



```
In [12]: # Correlation between features:
plt.figure(figsize=(10,10))
corrmat = wdata.corr(method='spearman')
sns.heatmap(corrmat)
# Findings
# There isn't any strong correlation between features and target variable.
# Alcohol shows moderate correlation.
# Density, Chlorides, volatile acidity show negative correlation which we also
# interpreted from inferences made using boxplot
```

Out[12]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7fad70ac7390>



```
In [69]: # Calculate the correlation values
feature_cols = wdata.columns[:-1]
corr_values = wdata[feature_cols].corr()


# Simplify by emptying all the data below the diagonal
tril_index = np.tril_indices_from(corr_values)

# Make the unused values NaNs
for coord in zip(*tril_index):
    corr_values.iloc[coord[0], coord[1]] = np.NaN

# Stack the data and convert to a data frame
corr_values = (corr_values
               .stack()
               .to_frame()
               .reset_index()
               .rename(columns={'level_0': 'feature1',
                               'level_1': 'feature2',
                               0: 'correlation'}))

# Get the absolute values for sorting
corr_values['abs_correlation'] = corr_values.correlation.abs()
```

```
In [70]: sns.set_context('talk')
sns.set_style('white')
ax = corr_values.abs_correlation.hist(bins=50, figsize=(12, 8))
ax.set(xlabel='Absolute Correlation', ylabel='Frequency');
```



```
In [71]: # The most highly correlated values ->> greater than 0.5
corr_values.sort_values('correlation', ascending=False).query('abs_correlation
>0.5')
### END SOLUTION
```

Out[71]:

	feature1	feature2	correlation	abs_correlation
45	free_sulfur_dioxide	total_sulfur_dioxide	0.720934	0.720934
33	residual_sugar	density	0.552517	0.552517
58	density	alcohol	-0.686745	0.686745

## C-2 : DATA CLEANING

```
In [13]: # Let's check dfor the null values in dataset
print(wdata.isnull().sum().to_string())
### there are no null values ###
```

fixed_acidity	0
volatile_acidity	0
citric_acid	0
residual_sugar	0
chlorides	0
free_sulfur_dioxide	0
total_sulfur_dioxide	0
density	0
pH	0
sulphates	0
alcohol	0
quality	0
color	0

## C-3 : FEATURE ENGINEERING

### Convert to a Classification Problem

Since the objective is to develop and evaluate various classification models, so I needed to change the output variable to a binary output. For this problem, I defined a bottle of wine as 'good quality' if it had a quality score of 7 or higher, and if it had a score of less than 7, it was deemed 'bad quality'.

Each wine in this dataset is given a “quality” score between 0 and 10. For the purpose of this project, I converted the output to a binary output where each wine is either “good quality” (a score of 7 or higher) or not (a score below 7). The quality of a wine is determined by 11 input variables: These datasets can be viewed as classification or regression tasks. The classes are ordered and not balanced (e.g. there are many more normal wines than excellent or poor ones). Outlier detection algorithms could be used to detect the few excellent or poor wines. Also, we are not sure if all input variables are relevant. So it could be interesting to test feature selection methods.

```
In [197]: # coverting y (color column) to integer 0,1
# Scikit Learn classifiers won't accept a sparse matrix for the prediction column. Thus, either LabelEncoder needs to be used to convert to integers,
# or if DictVectorizer is used, the resulting matrix must be converted to a non-sparse array.
# In this case since we had only two labels, we use simple astype(int) code as below

from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
wdatat=wdata.copy()
wdatat['color'] = le.fit_transform(wdatat.color)
wdatat.sample(5)
```

Out[197]:

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	free_sulfur_dioxide	tot
1614	6.6	0.17	0.38	1.5	0.032		28.0
808	7.4	0.53	0.12	1.9	0.165		4.0
3560	9.5	0.21	0.47	1.3	0.039		21.0
6287	6.7	0.16	0.32	12.5	0.035		18.0
3247	7.1	0.30	0.49	1.6	0.045		31.0

```
In [198]: wdatat['color'].value_counts()
```

```
Out[198]: 1    4898
0    1599
Name: color, dtype: int64
```

```
In [199]: # Each wine in this dataset is given a "quality" score between 0 and 10.
# For the purpose of this project, I converted the output to a binary output w
here
# Each wine is either "good quality" (a score of 7 or higher) or not (a score
below 7).
# Create Classification version of target variable
# wdata1 = wdata.copy(deep=True)
wdata1['goodquality'] = [1 if x >= 7 else 0 for x in wdata['quality']]

# DROPPING 'QUALITY' FEATURE AS ITS REPLACED BY 'GOODQUALITY'

wdata1 = wdata1.drop('quality', axis=1)
```

```
In [200]: wdata1.sample(5)
```

Out[200]:

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	free_sulfur_dioxide	tot
5467	6.0	0.220	0.25	11.1	0.056		112.0
2617	7.1	0.180	0.42	1.4	0.045		47.0
877	7.7	0.715	0.01	2.1	0.064		31.0
5837	6.4	0.290	0.18	15.0	0.040		21.0
5384	5.6	0.190	0.31	2.7	0.027		11.0

```
In [201]: ### check the balace in goodquality
wdata1['goodquality'].value_counts()
```

Out[201]: 0 5220  
1 1277  
Name: goodquality, dtype: int64

```
In [202]: wdata1.describe()
```

Out[202]:

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	free_sulfur_dioxid
count	6497.000000	6497.000000	6497.000000	6497.000000	6497.000000	6497.000000
mean	7.215307	0.339666	0.318633	5.443235	0.056034	30.52531
std	1.296434	0.164636	0.145318	4.757804	0.035034	17.74940
min	3.800000	0.080000	0.000000	0.600000	0.009000	1.000000
25%	6.400000	0.230000	0.250000	1.800000	0.038000	17.00000
50%	7.000000	0.290000	0.310000	3.000000	0.047000	29.00000
75%	7.700000	0.400000	0.390000	8.100000	0.065000	41.00000
max	15.900000	1.580000	1.660000	65.800000	0.611000	289.00000



```
In [203]: # Separate feature variables and target variable
X = wdatat.drop(['goodquality'], axis = 1)
y = wdatat['goodquality']
```

```
In [171]: ### SCALING OPTION1 ###
# Normalize feature variables
from sklearn.preprocessing import StandardScaler
#X_features = X
#X = StandardScaler().fit_transform(X)
```

```
In [204]: ### SCALING OPTION2 ###
from sklearn.preprocessing import MinMaxScaler
scalar = MinMaxScaler()
X = scalar.fit_transform(X)
```

```
In [205]: pd.DataFrame(X).describe()
```

Out[205]:

	0	1	2	3	4	5	
<b>count</b>	6497.000000	6497.000000	6497.000000	6497.000000	6497.000000	6497.000000	6497.000000
<b>mean</b>	0.282257	0.173111	0.191948	0.074283	0.078129	0.102518	0.252800
<b>std</b>	0.107143	0.109758	0.087541	0.072972	0.058195	0.061630	0.130200
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	0.214876	0.100000	0.150602	0.018405	0.048173	0.055556	0.163500
<b>50%</b>	0.264463	0.140000	0.186747	0.036810	0.063123	0.097222	0.258000
<b>75%</b>	0.322314	0.213333	0.234940	0.115031	0.093023	0.138889	0.345600
<b>max</b>	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

## SPLITTING THE DATA

we will Split the data into train and test data sets. This can be done using any method, but consider using Scikit-learn's StratifiedShuffleSplit to maintain the same ratio of predictor classes. Regardless of methods used to split the data, compare the ratio of classes in both the train and test splits.

```
In [173]: ### OPTION 1 SPLITTING - Normal Split ###
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.3, random_state=50)
```

```
In [222]: ### OPTION 2 SPLITTING - StratifiedShuffleSplit ###
from sklearn.model_selection import StratifiedShuffleSplit
feature_cols = [x for x in wdatat.columns if x not in 'goodquality']
# Split the data into two parts with 1000 points in the test data
# This creates a generator
strat_shuff_split = StratifiedShuffleSplit(n_splits=1, test_size=0.3, random_s
tate=50)

# Get the index values from the generator

train_idx, test_idx = next(strat_shuff_split.split(wdatat[feature_cols], wdata
t['goodquality']))

# Create the data sets
X_train = wdatat.loc[train_idx, feature_cols]
y_train = wdatat.loc[train_idx, 'goodquality']

X_test = wdatat.loc[test_idx, feature_cols]
y_test = wdatat.loc[test_idx, 'goodquality']
```

```
In [223]: wdatat['goodquality'].value_counts(normalize=True)
```

```
Out[223]: 0    0.803448
          1    0.196552
          Name: goodquality, dtype: float64
```

```
In [225]: y_train.value_counts(normalize=True)
```

```
Out[225]: 0    0.803387
          1    0.196613
          Name: goodquality, dtype: float64
```

```
In [224]: y_test.value_counts(normalize=True)
```

```
Out[224]: 0    0.80359
          1    0.19641
          Name: goodquality, dtype: float64
```

**D) Summary of training at least three different classifier models, preferably of different nature in explainability and predictability. For example, you can start with a simple logistic regression as a baseline, adding other models or ensemble models. Preferably, all your models use the same training and test splits, or the same cross-validation me**

We will now fit a logistic regression model without any regularization using all of the features. We will also use cross validation to determine the hyperparameters, fit models using L1, and L2 regularization and compare/evaluate each of these models as well.

MODEL 1 >>> Standard Logistic Regression, L1 Regularized & L2 Regularized

MODEL 2 >>> Support Vector Machine (SVM)

MODEL 3 >>> Decision Tree & GridSearchCV

## MODEL 1 >>> Standard Logistic Regression, L1 Regularized & L2 Regularized

```
In [226]: # Standard logistic regression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, roc_auc_score
lr = LogisticRegression(solver='liblinear').fit(X_train, y_train)
y_test_p_lr=lr.predict(X_test)
y_train_p_lr=lr.predict(X_train)
lr.coef_
```

```
Out[226]: array([[ -0.02308799, -3.38401235, -0.03411285,  0.04823298, -1.6922069 ,
  0.01330668, -0.00536188, -4.04852663, -0.65129519,  1.30774428,
  0.87766094, -0.19464785]])
```

```
In [105]: confusion_matrix(y_test, y_test_p_lr)
```

```
Out[105]: array([[1477,  62],
 [ 322,  89]])
```

```
In [106]: accuracy_score(y_test, y_test_p_lr)
```

```
Out[106]: 0.803076923076923
```

```
In [107]: # The error on the training and test data sets
train_test_lr_error = pd.concat([measure_error(y_train, y_train_p_lr, 'train'
),
                                measure_error(y_test, y_test_p_lr, 'test')],
                                axis=1)

train_test_lr_error
```

```
Out[107]:
```

	train	test
<b>accuracy</b>	0.821861	0.803077
<b>precision</b>	0.594595	0.589404
<b>recall</b>	0.203233	0.216545
<b>f1</b>	0.302926	0.316726

```
In [108]: print(classification_report(y_train, y_train_p_lr))
```

	precision	recall	f1-score	support
0	0.84	0.97	0.90	3681
1	0.59	0.20	0.30	866
accuracy			0.82	4547
macro avg	0.72	0.59	0.60	4547
weighted avg	0.79	0.82	0.78	4547

```
In [227]: print(classification_report(y_test, y_test_p_lr))
```

	precision	recall	f1-score	support
0	0.84	0.96	0.89	1567
1	0.56	0.23	0.33	383
accuracy			0.81	1950
macro avg	0.70	0.59	0.61	1950
weighted avg	0.78	0.81	0.78	1950

```
In [228]: # L1 regularized logistic regression
from sklearn.linear_model import LogisticRegressionCV
lrl1 = LogisticRegressionCV(Cs=10, cv=4, penalty='l1', solver='liblinear').fit(
    X_train, y_train)
y_test_p_lrl1=lrl1.predict(X_test)
y_train_p_lrl1=lrl1.predict(X_train)
lrl1.coef_
```

```
Out[228]: array([[ 1.49435344e-01, -3.67815311e+00, -2.08638263e-01,
                   6.69010295e-02, -9.97108097e+00,  1.39583730e-02,
                   -5.15169951e-03, -7.24327084e+00,  7.60464520e-01,
                   1.69965980e+00,  9.37068926e-01, -7.52412210e-02]])
```

```
In [111]: confusion_matrix(y_test, y_test_p_lrl1)
```

```
Out[111]: array([[1475,  64],
                  [ 319,  92]])
```

```
In [112]: accuracy_score(y_test, y_test_p_lrl1)
```

```
Out[112]: 0.8035897435897436
```

```
In [94]: # The error on the training and test data sets
train_test_lrl1_error = pd.concat([measure_error(y_train, y_train_p_lrl1, 'train'),
                                   measure_error(y_test, y_test_p_lrl1, 'test')],
                                   axis=1)

train_test_lrl1_error
```

Out[94]:

	train	test
<b>accuracy</b>	0.822740	0.808205
<b>precision</b>	0.578534	0.603352
<b>recall</b>	0.255196	0.262774
<b>f1</b>	0.354167	0.366102

```
In [116]: print(classification_report(y_train, y_train_p_lrl1))
```

	precision	recall	f1-score	support
0	0.84	0.96	0.90	3681
1	0.58	0.21	0.31	866
accuracy			0.82	4547
macro avg	0.71	0.59	0.60	4547
weighted avg	0.79	0.82	0.78	4547

```
In [229]: print(classification_report(y_test, y_test_p_lrl1))
```

	precision	recall	f1-score	support
0	0.84	0.95	0.89	1567
1	0.55	0.27	0.36	383
accuracy			0.81	1950
macro avg	0.70	0.61	0.63	1950
weighted avg	0.78	0.81	0.79	1950

```
In [230]: # L2 regularized logistic regression
from sklearn.linear_model import LogisticRegressionCV
lrl2 = LogisticRegressionCV(Cs=10, cv=4, penalty='l2', solver='liblinear').fit(
    X_train, y_train)
y_test_p_lrl2=lrl2.predict(X_test)
y_train_p_lrl2=lrl2.predict(X_train)
lrl2.coef_
```

```
Out[230]: array([[ 1.61722515e-01, -3.70870366e+00, -2.43352240e-01,
                   6.84485714e-02, -8.41867258e+00,  1.40428122e-02,
                   -5.21191047e-03, -7.90292366e+00,  8.98749353e-01,
                   1.69727696e+00,  9.49093785e-01,  1.28591709e-02]])
```

```
In [81]: # The error on the training and test data sets

train_test_lrl2_error = pd.concat([measure_error(y_train, y_train_p_lrl2, 'train'),
                                   measure_error(y_test, y_test_p_lrl2, 'test')],
                                   axis=1)

train_test_lrl2_error
```

Out[81]:

	train	test
<b>accuracy</b>	0.822520	0.809231
<b>precision</b>	0.576227	0.610169
<b>recall</b>	0.257506	0.262774
<b>f1</b>	0.355946	0.367347

```
In [34]: confusion_matrix(y_test, y_test_p_lrl2)
```

Out[34]: array([[1470, 69],  
[ 303, 108]])

```
In [35]: accuracy_score(y_test, y_test_p_lrl2)
```

Out[35]: 0.8092307692307692

```
In [89]: print(classification_report(y_train, y_train_p_lrl2))
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	3681
1	0.58	0.26	0.36	866
accuracy			0.82	4547
macro avg	0.71	0.61	0.63	4547
weighted avg	0.79	0.82	0.79	4547

```
In [231]: print(classification_report(y_test, y_test_p_lrl2))
```

	precision	recall	f1-score	support
0	0.84	0.95	0.89	1567
1	0.55	0.27	0.36	383
accuracy			0.81	1950
macro avg	0.69	0.61	0.62	1950
weighted avg	0.78	0.81	0.79	1950

For each model, calculate the following error metrics: Accuracy Precision Recall F-score Confusion Matrix  
Decide how to combine the multi-class metrics into a single value for each model.

## MODEL 2 >>> Support Vector Machine (SVM)

In [ ]:

```
In [232]: from sklearn.svm import LinearSVC
lscv = LinearSVC()
lscv.fit(X_train, y_train)
```

```
/opt/conda/envs/Python-3.7-OpenCE/lib/python3.7/site-packages/sklearn/svm/_base.py:977: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
```

Out[232]: LinearSVC()

```
In [119]: lscv.coef_
```

```
Out[119]: array([[ 0.01042635, -1.240539, -0.27502932,  0.02136014, -0.62806316,
                  0.01712583, -0.00208026, -1.38674335, -0.31036954,  0.88079438,
                  0.29087225, -0.1104003 ]])
```

```
In [235]: y_test_p_LSCV = lscv.predict(X_test)
y_train_p_LSCV = lscv.predict(X_train)
confusion_matrix(y_test, y_test_p_LSCV)
```

```
Out[235]: array([[1560,  7],
                 [ 374,  9]])
```

```
In [121]: ### The decision tree predicts a little better on the training data than the test data, which is consistent with (mild) overfitting. Also notice the perfect recall score for the training data. In many instances, this prediction difference is even greater than that seen here.
# The error on the training and test data sets
train_test_lscv_error = pd.concat([measure_error(y_train, y_train_p_LSCV, 'train'),
                                   measure_error(y_test, y_test_p_LSCV, 'test')],
                                   axis=1)

train_test_lscv_error
```

Out[121]:

	train	test
<b>accuracy</b>	0.768199	0.778462
<b>precision</b>	0.406931	0.475177
<b>recall</b>	0.474596	0.489051
<b>f1</b>	0.438166	0.482014

```
In [122]: accuracy_score(y_test, y_test_p_LSCV)
```

Out[122]: 0.7784615384615384

```
In [149]: print(classification_report(y_train, y_train_p_LSCV))
```

	precision	recall	f1-score	support
0	0.87	0.84	0.85	3681
1	0.41	0.47	0.44	866
accuracy			0.77	4547
macro avg	0.64	0.66	0.65	4547
weighted avg	0.78	0.77	0.77	4547

```
In [236]: print(classification_report(y_test, y_test_p_LSCV))
```

	precision	recall	f1-score	support
0	0.81	1.00	0.89	1567
1	0.56	0.02	0.05	383
accuracy			0.80	1950
macro avg	0.68	0.51	0.47	1950
weighted avg	0.76	0.80	0.72	1950

## MODEL 3 >>> Decision Tree & GridSearchCV



## 1- Decision Tree - DecisionTreeClassifier

```
In [45]: from sklearn.metrics import classification_report
```

```
In [237]: from sklearn.tree import DecisionTreeClassifier
dtc = DecisionTreeClassifier(random_state=1)
dtc.fit(X_train, y_train)
dtc.tree_.node_count, dtc.tree_.max_depth
```

Out[237]: (1059, 25)

```
In [152]: ### The decision tree predicts a little better on the training data than the t
est data, which is consistent with (mild) overfitting. Also notice the perfect
recall score for the training data. In many instances, this prediction differe
nce is even greater than that seen here.
# The error on the training and test data sets
train_test_dtc_error = pd.concat([measure_error(y_train, y_train_p_dtc, 'train'),
                                measure_error(y_test, y_test_p_dtc, 'test')],
                                axis=1)

train_test_dtc_error
```

Out[152]:

	train	test
<b>accuracy</b>	1.0	0.826154
<b>precision</b>	1.0	0.588670
<b>recall</b>	1.0	0.581509
<b>f1</b>	1.0	0.585067

```
In [239]: y_test_p_dtc = dtc.predict(X_test)
y_train_p_dtc = dtc.predict(X_train)
print(classification_report(y_test, y_test_p_dtc))
```

	precision	recall	f1-score	support
0	0.91	0.90	0.90	1567
1	0.60	0.62	0.61	383
accuracy			0.84	1950
macro avg	0.75	0.76	0.76	1950
weighted avg	0.85	0.84	0.84	1950

```
In [155]: print(classification_report(y_train, y_train_p_dtc))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3681
1	1.00	1.00	1.00	866
accuracy			1.00	4547
macro avg	1.00	1.00	1.00	4547
weighted avg	1.00	1.00	1.00	4547

```
In [156]: confusion_matrix(y_test, y_test_p_dtc)
```

```
Out[156]: array([[1372, 167],
                [ 172, 239]])
```

```
In [131]: accuracy_score(y_test, y_test_p_dtc)
```

```
Out[131]: 0.8261538461538461
```

## 2- Decision Tree - Random Forest

```
In [240]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(random_state=1)
rfc.fit(X_train, y_train)
```

```
Out[240]: RandomForestClassifier(random_state=1)
```

```
In [133]: ### The decision tree predicts a little better on the training data than the t
est data, which is consistent with (mild) overfitting. Also notice the perfect
recall score for the training data. In many instances, this prediction differe
nce is even greater than that seen here.
# The error on the training and test data sets
train_test_rfc_error = pd.concat([measure_error(y_train, y_train_p_rfc, 'train'),
                                   measure_error(y_test, y_test_p_rfc, 'test')],
                                   axis=1)

train_test_rfc_error
```

```
Out[133]:
```

	train	test
accuracy	1.0	0.867179
precision	1.0	0.781481
recall	1.0	0.513382
f1	1.0	0.619677

```
In [134]: confusion_matrix(y_test, y_test_p_rfc)
```

```
Out[134]: array([[1480,   59],
                 [ 200,  211]])
```

```
In [137]: accuracy_score(y_test, y_test_p_rfc)
```

```
Out[137]: 0.8676923076923077
```

```
In [241]: y_test_p_rfc = rfc.predict(X_test)
          print(classification_report(y_test, y_test_p_rfc))
```

	precision	recall	f1-score	support
0	0.90	0.97	0.93	1567
1	0.80	0.56	0.66	383
accuracy			0.89	1950
macro avg	0.85	0.76	0.80	1950
weighted avg	0.88	0.89	0.88	1950

```
In [139]: y_train_p_rfc = rfc.predict(X_train)
          print(classification_report(y_train, y_train_p_rfc))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3681
1	1.00	1.00	1.00	866
accuracy			1.00	4547
macro avg	1.00	1.00	1.00	4547
weighted avg	1.00	1.00	1.00	4547

### 3- Decision Tree - GridSearchCV

Using grid search with cross validation, find a decision tree that performs well on the test data set.

```
In [242]: from sklearn.model_selection import GridSearchCV

param_grid = {'max_depth':range(1, dtc.tree_.max_depth+1, 2),
              'max_features': range(1, len(dtc.feature_importances_)+1)}

gscv = GridSearchCV(DecisionTreeClassifier(random_state=42),
                    param_grid=param_grid,
                    scoring='accuracy',
                    n_jobs=-1)

gscv= gscv.fit(X_train, y_train)
y_train_p_gscv = gscv.predict(X_train)
y_test_p_gscv = gscv.predict(X_test)
```

```
In [192]: gscv.best_estimator_.tree_.node_count, gscv.best_estimator_.tree_.max_depth
```

```
Out[192]: (893, 13)
```

```
In [142]: ### TRAIN VS TEST ERROR - A function to return error metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

def measure_error(y_true, y_pred, label):
    return pd.Series({'accuracy':accuracy_score(y_true, y_pred),
                     'precision': precision_score(y_true, y_pred),
                     'recall': recall_score(y_true, y_pred),
                     'f1': f1_score(y_true, y_pred)},
                    name=label)
```

```
In [143]: ### The decision tree predicts a little better on the training data than the t
est data, which is consistent with (mild) overfitting. Also notice the perfect
recall score for the training data. In many instances, this prediction differe
nce is even greater than that seen here.
# The error on the training and test data sets
train_test_gscv_error = pd.concat([measure_error(y_train, y_train_p_gscv, 'tra
in'),
                                   measure_error(y_test, y_test_p_gscv, 'test')],
                                   axis=1)

train_test_gscv_error
```

```
Out[143]:
```

	train	test
<b>accuracy</b>	0.981526	0.830256
<b>precision</b>	0.942308	0.597087
<b>recall</b>	0.961894	0.598540
<b>f1</b>	0.952000	0.597813

```
In [144]: confusion_matrix(y_test, y_test_p_gscv)
```

```
Out[144]: array([[1373, 166],
                  [ 165, 246]])
```

```
In [145]: accuracy_score(y_test, y_test_p_gscv)
```

```
Out[145]: 0.8302564102564103
```

```
In [146]: print(classification_report(y_train, y_train_p_gscv))
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	3681
1	0.94	0.96	0.95	866
accuracy			0.98	4547
macro avg	0.97	0.97	0.97	4547
weighted avg	0.98	0.98	0.98	4547

```
In [243]: print(classification_report(y_test, y_test_p_gscv))
```

	precision	recall	f1-score	support
0	0.91	0.90	0.90	1567
1	0.60	0.63	0.61	383
accuracy			0.85	1950
macro avg	0.75	0.76	0.76	1950
weighted avg	0.85	0.85	0.85	1950

**E) A paragraph explaining which of your classifier models you recommend as a final model that best fits your needs in terms of accuracy and explainability.**

The objective of the analysis was to predict the quality of wine. The accuracy of the model has to be very high. At the same time since the intent is to reject low quality wines and select good quality wines, it's important to have good recall for good quality wines. I will compare my models based on these two metrics (overall accuracy and recall) and select the final model.

Here is summary of metrics for various models

Model Accuracy Recall

Logistic Regression (3 variations) 0.80 0.22

Logit 0.80 0.22

Logit2 0.81 0.26

Support Vector Machine svm 0.78 0.49

Decision Tree (3 variations)

dct 0.83 0.58

rfc 0.87 0.52

gscv 0.83 0.60

Overall Decision tree is giving best performance. Further using splitshuffle splitting we further enhance its performance

Model Accuracy Recall

dct 0.84 0.62

rfc 0.89 0.56

gscv 0.85 0.63

Finally taking into account both metrics, I finalize to use RANDOM FOREST for my prediction.

**F) Summary Key Findings and Insights, which walks your reader through the main drivers of your model and insights from your data derived from your classifier model.**

The objective of the analysis was to predict the quality of wine. The accuracy of the model has to be very high. At the same time since the intent is to reject low quality wines and select good quality wines, it's important to have good recall for good quality wines. I evaluated my models based on these two metrics (overall accuracy and recall) and select the final model.

While the accuracy of the model is quite good, recall is very poor specifically in Logistic regression. If the objective of the model is very much dependent on recall, it's worthwhile to explore why it is so low. The best recall value I could get in 63% which is not very satisfactory.

## **G) Suggestions for next steps in analyzing this data, which may include suggesting revisiting this model after adding specific data features that may help you achieve a better explanation or a better prediction.**

There are other models which could have been tried but for the purpose of this assignment I just tried these three models

To further enhance accuracy/recall we could have used multiple bagging and boosting techniques

I spent little time on feature engineering and applied only few techniques. Time permitting, I could have done more feature engineering like elimination of few features which are not relevant based on data exploration I have done.

Overall despite achieving good overall accuracy (89%) there are other tools & techniques which could have further improved my model.

In [ ]: